

AdaBoost

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON



Elie Kawerk

Data Scientist

Boosting

- **Boosting:** Ensemble method combining several weak learners to form a strong learner.
- **Weak learner:** Model doing slightly better than random guessing.
- Example of weak learner: Decision stump (CART whose maximum depth is 1).


Boosting

- Train an ensemble of predictors sequentially.
- Each predictor tries to correct its predecessor.
- Most popular boosting methods:
 - AdaBoost,
 - Gradient Boosting.

Adaboost


- Stands for **Adaptive Boosting**.
- Each predictor pays more attention to the instances wrongly predicted by its predecessor.
- Achieved by changing the weights of training instances.
- Each predictor is assigned a coefficient α .
- α depends on the predictor's training error.

AdaBoost: Training



Learning Rate

Learning rate: $0 < \eta \leq 1$



AdaBoost: Prediction

- Classification:
 - Weighted majority voting.
 - In sklearn: `AdaBoostClassifier`.
- Regression:
 - Weighted average.
 - In sklearn: `AdaBoostRegressor`.

AdaBoost Classification in sklearn (Breast Cancer dataset)

```
# Import models and utility functions
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split

# Set seed for reproducibility
SEED = 1

# Split data into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                    stratify=y,
                                                    random_state=SEED)
```

```
# Instantiate a classification-tree 'dt'  
dt = DecisionTreeClassifier(max_depth=1, random_state=SEED)  
  
# Instantiate an AdaBoost classifier 'adb_clf'  
adb_clf = AdaBoostClassifier(base_estimator=dt, n_estimators=100)  
  
# Fit 'adb_clf' to the training set  
adb_clf.fit(X_train, y_train)  
  
# Predict the test set probabilities of positive class  
y_pred_proba = adb_clf.predict_proba(X_test)[:,1]  
  
# Evaluate test-set roc_auc_score  
adb_clf_roc_auc_score = roc_auc_score(y_test, y_pred_proba)
```

AdaBoost Classification in sklearn (Breast Cancer dataset)

```
# Print adb_clf_roc_auc_score  
print('ROC AUC score: {:.2f}'.format(adb_clf_roc_auc_score))
```

```
ROC AUC score: 0.99
```

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Gradient Boosting (GB)

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON




Elie Kawerk

Data Scientist


Gradient Boosted Trees

- Sequential correction of predecessor's errors.
- Does not tweak the weights of training instances.
- Fit each predictor is trained using its predecessor's residual errors as labels.
- Gradient Boosted Trees: a CART is used as a base learner.

Gradient Boosted Trees for Regression: Training



Shrinkage



Gradient Boosted Trees: Prediction

- Regression:
 - $y_{pred} = y_1 + \eta r_1 + \dots + \eta r_N$
 - In sklearn: `GradientBoostingRegressor` .
- Classification:
 - In sklearn: `GradientBoostingClassifier` .

Gradient Boosting in sklearn (auto dataset)

```
# Import models and utility functions
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as MSE

# Set seed for reproducibility
SEED = 1

# Split dataset into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(X,y,
                                                    test_size=0.3,
                                                    random_state=SEED)
```

```
# Instantiate a GradientBoostingRegressor 'gbt'  
gbt = GradientBoostingRegressor(n_estimators=300, max_depth=1, random_state=SEED)  
  
# Fit 'gbt' to the training set  
gbt.fit(X_train, y_train)  
  
# Predict the test set labels  
y_pred = gbt.predict(X_test)  
  
# Evaluate the test set RMSE  
rmse_test = MSE(y_test, y_pred)**(1/2)  
  
# Print the test set RMSE  
print('Test set RMSE: {:.2f}'.format(rmse_test))
```

Test set RMSE: 4.01

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON

Stochastic Gradient Boosting (SGB)

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON



Elie Kawerk

Data Scientist


Gradient Boosting: Cons

- GB involves an exhaustive search procedure.
- Each CART is trained to find the best split points and features.
- May lead to CARTs using the same split points and maybe the same features.

Stochastic Gradient Boosting

- Each tree is trained on a random subset of rows of the training data.
- The sampled instances (40%-80% of the training set) are sampled without replacement.
- Features are sampled (without replacement) when choosing split points.
- Result: further ensemble diversity.
- Effect: adding further variance to the ensemble of trees.

Stochastic Gradient Boosting: Training



Stochastic Gradient Boosting in sklearn (auto dataset)

```
# Import models and utility functions
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as MSE

# Set seed for reproducibility
SEED = 1

# Split dataset into 70% train and 30% test
X_train, X_test, y_train, y_test = train_test_split(X,y,
                                                    test_size=0.3,
                                                    random_state=SEED)
```

Stochastic Gradient Boosting in sklearn (auto dataset)

```
# Instantiate a stochastic GradientBoostingRegressor 'sgbt'  
sgbt = GradientBoostingRegressor(max_depth=1,  
                                  subsample=0.8,  
                                  max_features=0.2,  
                                  n_estimators=300,  
                                  random_state=SEED)  
  
# Fit 'sgbt' to the training set  
sgbt.fit(X_train, y_train)  
  
# Predict the test set labels  
y_pred = sgbt.predict(X_test)
```

Stochastic Gradient Boosting in sklearn (auto dataset)

```
# Evaluate test set RMSE 'rmse_test'  
rmse_test = MSE(y_test, y_pred)**(1/2)  
  
# Print 'rmse_test'  
print('Test set RMSE: {:.2f}'.format(rmse_test))
```

Test set RMSE: 3.95

Let's practice!

MACHINE LEARNING WITH TREE-BASED MODELS IN PYTHON